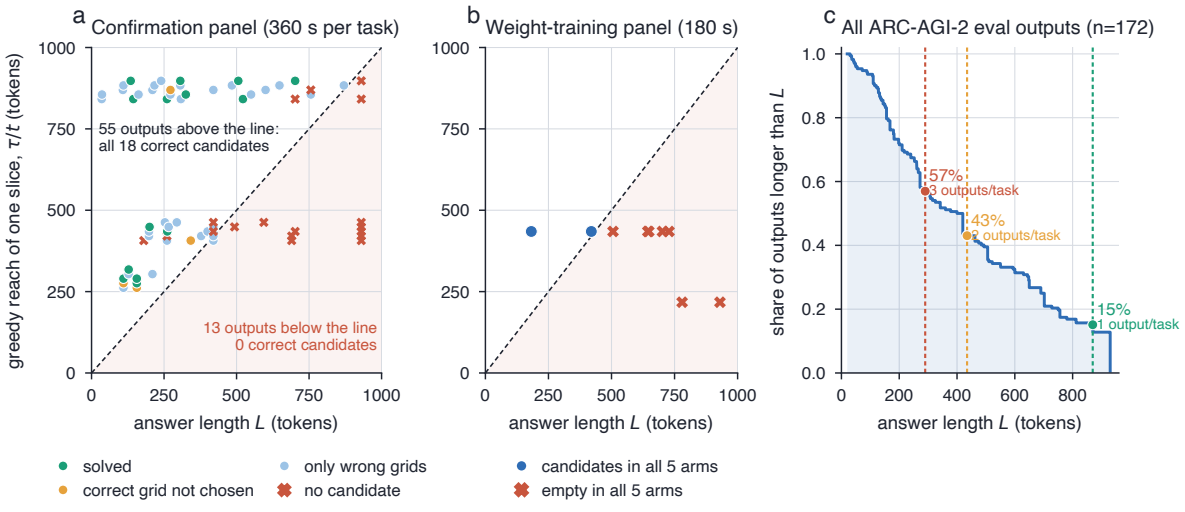
a
Confirmation panel (360 s per task)
55 outputs above the line:
all 18 correct candidates
13 outputs below the line
0 correct candidates
greedy reach of one slice, τ/t (tokens)
answer length L (tokens)
b
Weight-training panel (180 s)
answer length L (tokens)
c
All ARC-AGI-2 eval outputs (n=172)
share of outputs longer than L
57%
3 outputs/task
43%
2 outputs/task
15%
1 output/task
answer length L (tokens)
solved
correct grid not chosen
only wrong grids
no candidate
candidates in all 5 arms
empty in all 5 arms